

ML and NLP with Python

Python Libraries

- **NumPy** **Numerical computing, arrays**
- **Pandas** **Data manipulation**
- **Matplotlib** **Data visualization**
- **Seaborn** **Statistical data visualization**
- **Scikit-Learn** **Machine learning algorithms**
- TensorFlow Deep learning, neural networks
- Keras High-level API for deep learning
- PyTorch Deep learning (research-focused)
- XGBoost Gradient boosting for structured data
- LightGBM Fast boosting algorithm
- OpenCV Computer vision and image processing
- **NLTK** **Natural language processing**
- SpaCy Advanced NLP

scikit

- scikit-learn (sklearn) is a powerful machine learning library in Python that provides tools for:
 - Data Preprocessing (handling missing data, scaling, encoding)
 - Feature Extraction (Bag of Words, TF-IDF, PCA)
 - Supervised Learning (Regression & Classification models)
 - Unsupervised Learning (Clustering, Anomaly Detection)
 - Model Selection & Evaluation (Cross-validation, Hyperparameter tuning)

- Task 1: Load & Explore a Dataset

```
import pandas as pd
df = pd.read_csv('data.csv') # Load dataset
print(df.head()) # Show first 5 rows
print(df.info()) # Dataset summary
print(df.describe()) # Statistical summary
```

- Task 2: Train-Test Split

```
from sklearn.model_selection import train_test_split  
X = df.drop('Target', axis=1) # Features  
y = df['Target'] # Labels  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

- Task 3: Linear Regression

```
from sklearn.linear_model import LinearRegression  
model = LinearRegression()  
model.fit(X_train, y_train) # Train model  
y_pred = model.predict(X_test) # Make predictions
```

The `random_state` parameter ensures that the data split is reproducible. It controls the randomness of the train-test split, meaning:

Same `random_state` → Same Split Every Time

Different `random_state` → Different Split Every Time

- `from sklearn.model_selection import train_test_split`
- `import numpy as np`

- `# Create a simple dataset`
- `X = np.array(range(10)).reshape(-1, 1) # Features (0 to 9)`
- `y = np.array(range(10)) # Labels (0 to 9)`

- `# Split without random_state (results will change every time)`
- `X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)`

- `print("X_test:", X_test.ravel()) # Different results each time`

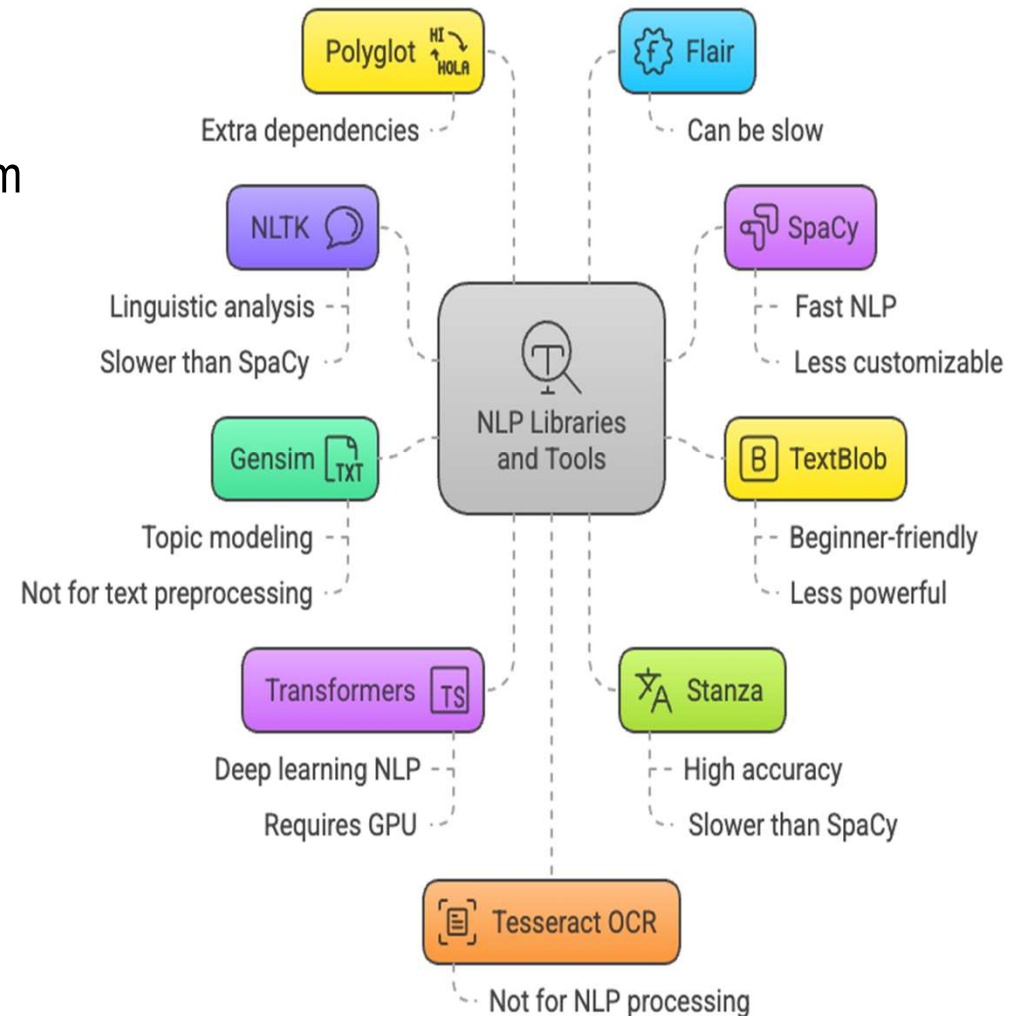
- Task 4: Logistic Regression

```
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
clf = LogisticRegression()
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

NLP Libraries in Python

- Python has number of libraries for NLP to perform tokenization, sentiment analysis, machine translation, text summarization, and more.
 - **NLTK (Natural Language Toolkit)**
 - **spaCy**
 - **TextBlob**
 - **Transformers (by Hugging Face)**
 - **Gensim**
 - **Tesseract OCR (for Text Extraction from Images)**
 - **Polyglot**
 - **Keras (for deep learning NLTK)**

NLP Libraries and Tools: Strengths and Weaknesses



NLP-II

BoW in Python

```
from sklearn.feature_extraction.text import CountVectorizer
texts = ["I love machine learning", "Machine learning is amazing", "I love coding"]
vectorizer = CountVectorizer()
bow = vectorizer.fit_transform(texts) //Learn the vocabulary dictionary and return document-term matrix.
print(vectorizer.get_feature_names_out())
print(bow.toarray())
```

```
['amazing' 'coding' 'is' 'learning' 'love' 'machine']
[[0 0 0 1 1 1]
 [1 0 1 1 0 1]
 [0 1 0 0 1 0]]
```

```
from sklearn.feature_extraction.text import CountVectorizer
>>> corpus = [ ... 'This is the first document.', ... 'This document is the second document.',
... 'And this is the third one.', ... 'Is this the first document?', ... ]
>>> vectorizer = CountVectorizer()
>>> X = vectorizer.fit_transform(corpus)
>>> vectorizer.get_feature_names_out()
array(['and', 'document', 'first', 'is', 'one', 'second', 'the', 'third', 'this'], ...)
>>> print(X.toarray())
[[0 1 1 1 0 0 1 0 1]
 [0 2 0 1 0 1 1 0 1]
 [1 0 0 1 1 0 1 1 1]
 [0 1 1 1 0 0 1 0 1]]
```

```
vectorizer2 = CountVectorizer(analyzer='word', ngram_range=(2, 2))
```

```
>>> X2 = vectorizer2.fit_transform(corpus)
```

```
>>> vectorizer2.get_feature_names_out()
```

```
array(['and this', 'document is', 'first document', 'is the', 'is this', 'second  
document', 'the first', 'the second', 'the third', 'third one', 'this document',  
'this is', 'this the'], ...)
```

```
>>> print(X2.toarray())
```

```
[[0 0 1 1 0 0 1 0 0 0 0 1 0]
```

```
[0 1 0 1 0 1 0 1 0 0 1 0 0]
```

```
[1 0 0 1 0 0 0 0 1 1 0 1 0]
```

```
[0 0 1 0 1 0 1 0 0 0 0 0 1]]
```

TF-IDF in Python

```
from sklearn.feature_extraction.text import TfidfVectorizer  
tfidf = TfidfVectorizer()  
X = tfidf.fit_transform(texts)  
print(tfidf.get_feature_names_out())  
print(X.toarray())
```

```
['amazing' 'coding' 'is' 'learning' 'love' 'machine']  
[[0. 0. 0. 0.57735027 0.57735027 0.57735027]  
 [0.5628291 0. 0.5628291 0.42804604 0. 0.42804604]  
 [0. 0.79596054 0. 0. 0.60534851 0. ]]
```

Similarity in Texts

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

```
text1 = set("machine learning is fun".split())  
text2 = set("learning about machine intelligence".split())  
jaccard = len(text1 & text2) / len(text1 | text2)  
print("Jaccard Similarity:", jaccard)
```

Jaccard Similarity: 0.3333333333333333

Cosine Similarity

$$\text{cosine similarity} = S_C(A, B) := \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|}$$

```
from sklearn.metrics.pairwise import cosine_similarity
tfidf_vec = TfidfVectorizer()
vecs = tfidf_vec.fit_transform(["machine learning is fun", "learning about machine intelligence"])
cos_sim = cosine_similarity(vecs[0:1], vecs[1:2])
print("Cosine Similarity:", cos_sim[0][0])
```

Cosine Similarity: 0.3360969272762575

Jaccard compares token sets;

Cosine compares vector angles (good for longer texts).

Sentiment Analysis

```
from textblob import TextBlob  
review = "The service was excellent and the staff was friendly."  
blob = TextBlob(review)  
print("Polarity:", blob.sentiment.polarity)  
print("Subjectivity:", blob.sentiment.subjectivity)
```


Word Cloud

```
from wordcloud import WordCloud
import matplotlib.pyplot as plt
text = "Python is simple and powerful. I love Python programming!"
wordcloud = WordCloud().generate(text)
plt.imshow(wordcloud, interpolation='bilinear')
plt.axis('off')
plt.show()
```

Text Generation using Keras

```
from keras.preprocessing.text import Tokenizer
from keras.utils import pad_sequences
text = "Machine learning is fun and exciting to learn"
tokenizer = Tokenizer()
tokenizer.fit_on_texts([text])
sequences = []
words = text.split()
for i in range(1, len(words)):
    seq = words[i+1]
    tokenized_seq = tokenizer.texts_to_sequences([' '.join(seq))][0]
    sequences.append(tokenized_seq)
# Pad the sequences
padded = pad_sequences(sequences)
print(padded)
```

Build a Model (LSTM Example)

```
from keras.models import Sequential
from keras.layers import Embedding, LSTM, Dense
model = Sequential()
model.add(Embedding(input_dim=50, output_dim=10,
input_length=padded.shape[1]))
model.add(LSTM(50))
model.add(Dense(50, activation='relu'))
model.add(Dense(len(tokenizer.word_index) + 1, activation='softmax'))
model.compile(loss='sparse_categorical_crossentropy', optimizer='adam')
# Normally you'd train the model with model.fit(), then use it to predict.
```

- **Long Short-Term Memory.**

It is a **type of Recurrent Neural Network (RNN)** that is specially designed to **remember long sequences** and patterns in data — especially useful in **Natural Language Processing (NLP)**, time series, and speech.